Work should never lose the reason it exists.

A compact view of infrastructure that keeps local action connected to purpose, authority, responsibility, effects, and evidence.

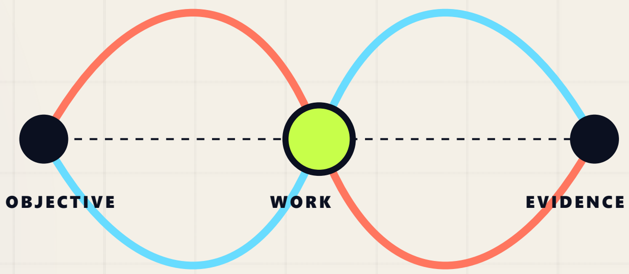

Classic editorial structure.
High-tech signal logic.
One inspectable line from intention to consequence.

An objective begins as a situated possibility.

Someone notices that reality could be different: a need, a problem, an opportunity, or a condition worth preserving.

Perception is never neutral or complete. Experience, knowledge, biology, social position, and culture influence what becomes visible and which distinctions seem important. These conditions do not mechanically determine the objective, but they shape the field from which it emerges.

The infrastructure does not need to model human perception. It needs to preserve the attributable transition from a situated perception to a declared objective - and allow that objective to be revised without losing its history.

> The system begins where possibility becomes an explicit, revisable intention.

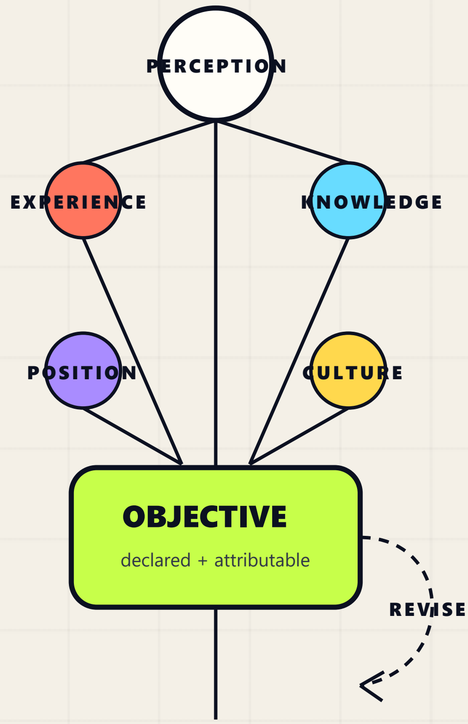

Definition is recursive, not merely sequential.

A monolithic intention gains contour as outcomes, constraints, questions, and routes become distinguishable. The parts compose the whole - and can change it.

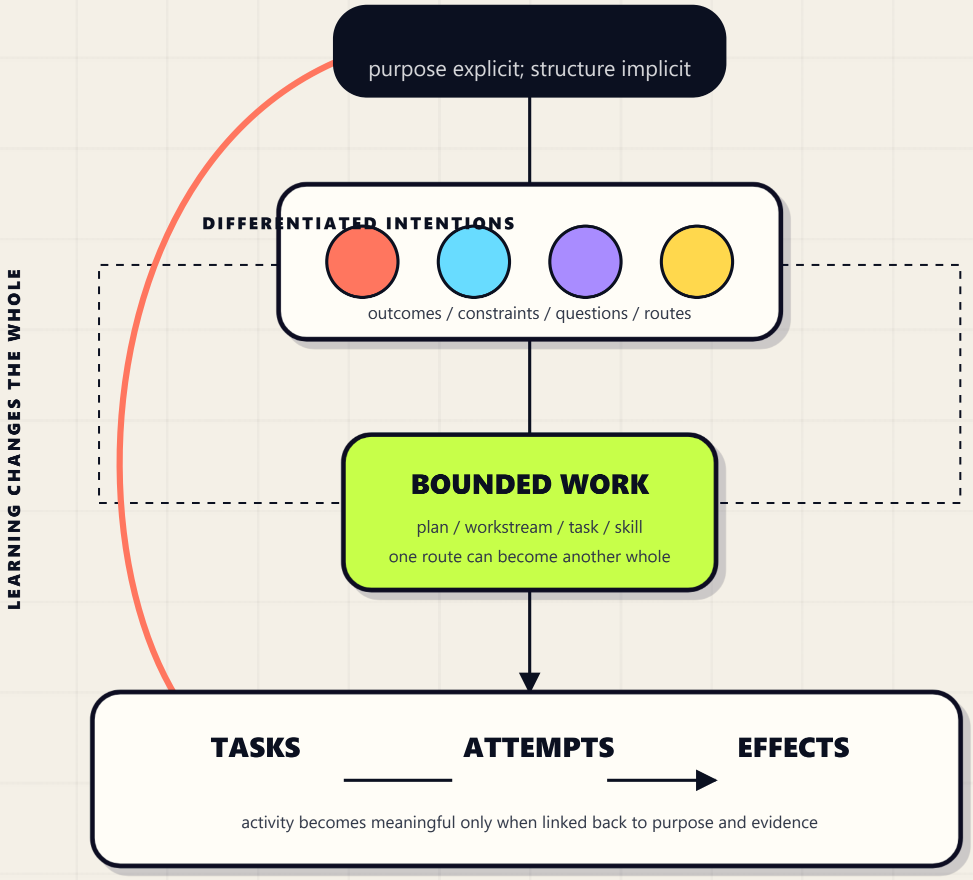

Not a mandatory lifecycle. A possible route may become one directly bounded task. The visualization shows decomposition possibilities, not required bureaucracy.

The visibility trade. Definition creates local precision. It also distributes meaning across conversations, plans, decisions, tasks, attempts, files, and effects.

Local correctness is not enough.

A task can be clear, assigned, and completed while the objective, assumption, or authority that once justified it has already changed.

> The failure is not missing activity. It is loss of context between levels of work.

Work-context infrastructure preserves the relationship between a local act and the larger contexts that explain it. It must let a person inspect different kinds of connection without treating one as proof of another.

01 Purpose	02 Authority	03 Assignment	04 Causation	05 Realization
Why does this work exist? Which objective does it serve?	What permits this actor or system to act?	Who carries responsibility for this bounded work?	Which attempt produced which concrete change?	What evidence supports a claimed outcome?

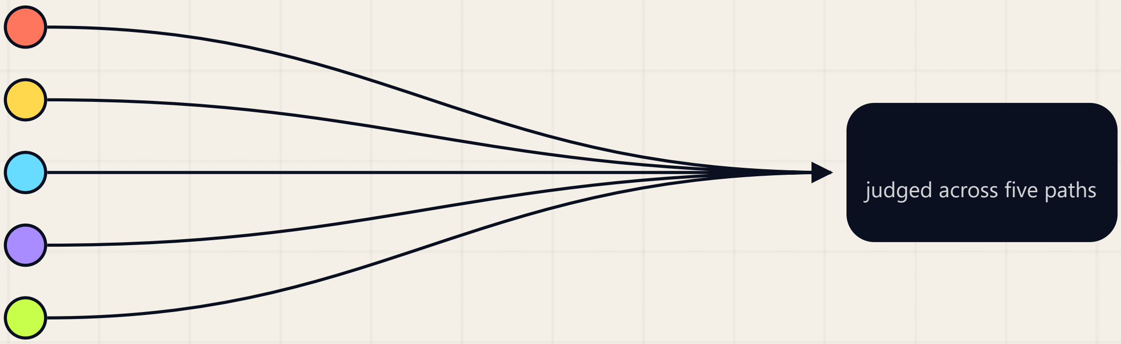

Independence matters. Assignment does not prove authority. Temporal order does not prove causation. An artifact does not prove realization. Successful execution does not prove purpose. A path may be witnessed, missing, conflicted, or superseded; the system must not invent a missing link.

Familiar systems contribute pieces, not the whole guarantee.

The claim is about preservation across the route from intention to effect - not the rejection of task tools, repositories, workflows, or orchestrators.

FRAMING	WHAT IT CONTRIBUTES	WHAT REMAINS UNPROVEN IF USED ALONE
Task management	Assignment, schedule, status.	Justification, authority, outcome evidence.
Document repository	Artifacts, versions, retrieval.	Typed relations and operational consequence.
Workflow engine	Repeatable transitions and execution.	Validity of meaning, authority, acceptance.
Multi-agent orchestrator	Dispatch, interaction, execution.	Lineage from intention to accepted effect.

MOVE UPWARD

From the part to the whole

From a task, attempt, artifact, or effect: recover the work that asked for it, the objective it serves, the decision that authorized it, and the evidence behind those choices.

Detects: orphan work, missing purpose, ambiguous authority.

MOVE DOWNWARD

From the whole to its parts

From an objective: inspect active routes, changed decisions, current attempts, produced effects, and evidence that the objective is advancing.

Detects: unrealized aspiration, incomplete routes, unverified implementation.

A

Keep links typed

A relationship says what it means. Proximity and chronology are not semantics.

B

Keep facts attributable

Every declaration, decision, attempt, correction, and effect has an accountable source.

C

Keep history appendable

Revision changes the current view without silently replacing the route that produced it.

PROPOSAL STATUS The scope comparison above is a hypothesis advanced by the source essay. Accepting it would define a problem boundary; it would not select an architecture or reject these systems as components.

Trust is composed across responsibility planes.

No service earns trust alone. Each boundary is justified by a failure it prevents and by the neighboring services required to complete a guarantee.

Entrance and decision plane

Turns bounded intent into an explicit epistemic strategy, a classified work kind, and a saved task graph that can be interpreted without silent mutation.

Subagents Strategy **Classifier** **Protocol** **Compiler**

Execution and mediation plane

Runs only the confirmed Dispatch. Fabric transports sealed exchanges without authoring, ranking, or summarizing their content.

Execution **Fabric** **CONFIRMATION GATE**

Knowledge plane

Supplies stable identities, typed relations, version-bound instructions, and the append-only record that makes every act recoverable.

Identity **Prompts** **Record**

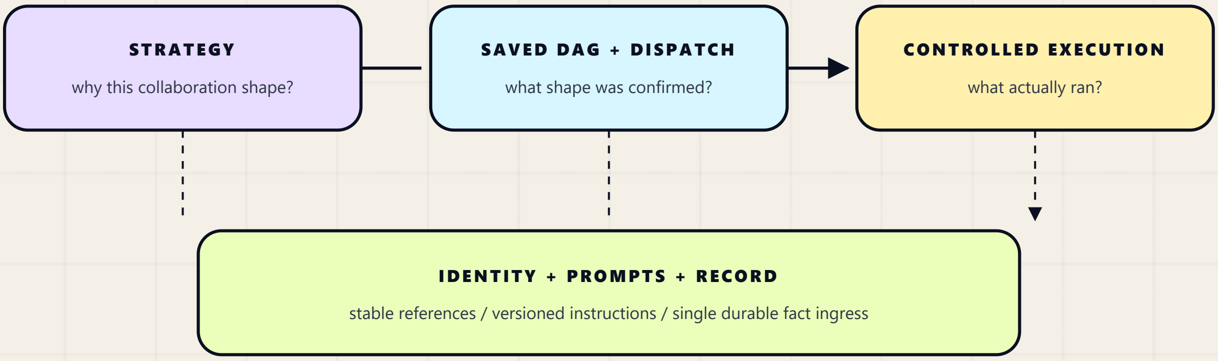

PROPOSED · NOT IMPLEMENTED **Architecture status.** The source baseline names eight services. This hypothesis adds Subagents Strategy as a ninth entrance service and refines Protocol into versioned per-task DAGs.

Separate the strategy, the saved graph, and the executable Dispatch.

The architecture preserves interpretation as evidence. Runtime must not smuggle semantic choices into execution.

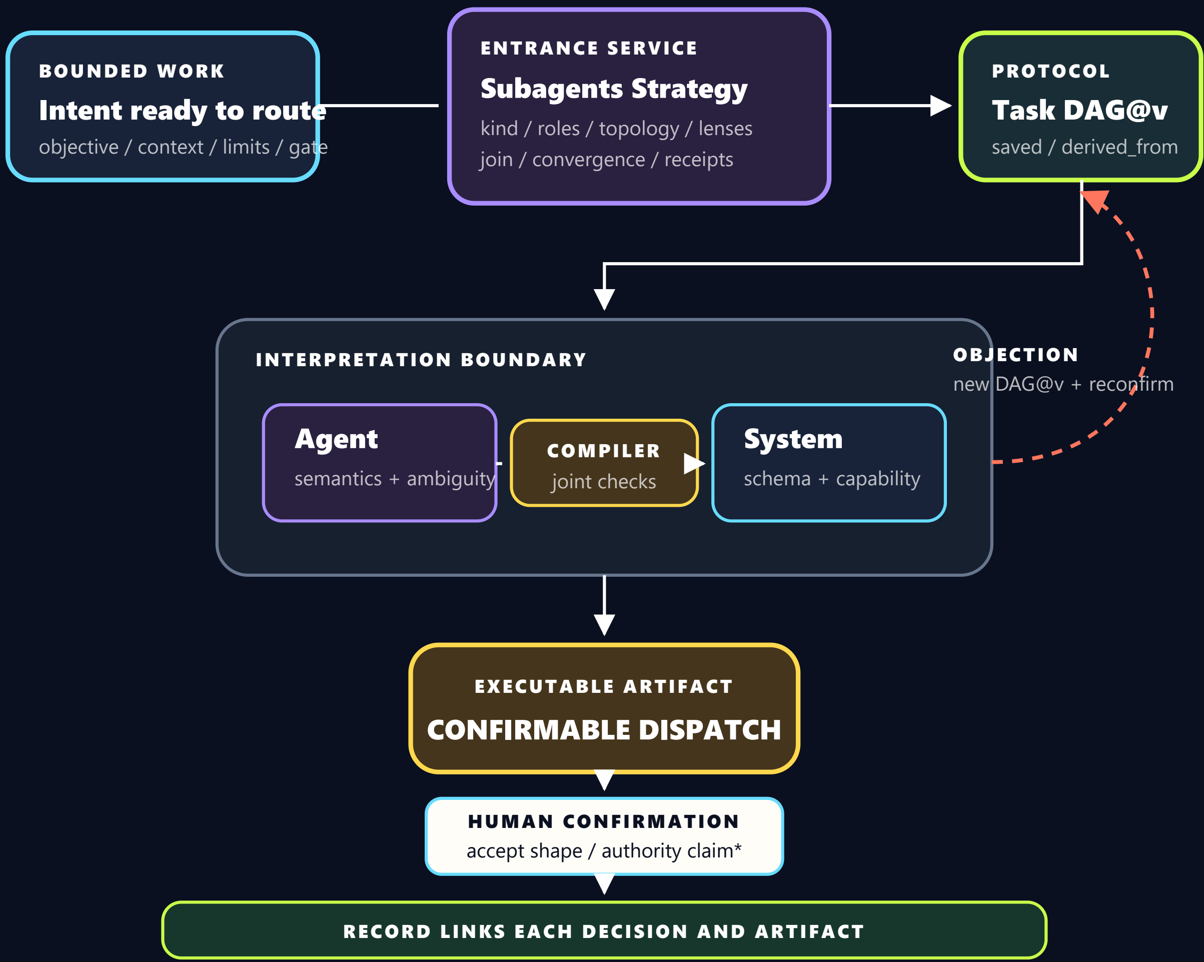

KEY DISTINCTION A DAG describes intended work. A Dispatch is its validated, bound, executable interpretation. A rejected interpretation appends a typed objection, produces a new DAG version, and reopens confirmation.

SERVICE MAP Strategy + Classifier -> Protocol DAG -> Compiler -> human confirmation -> Execution + Fabric. Identity and Prompts support the route; Record links its durable facts.

Execution should leave an explanation, not just an output.

The operational loop is complete only when the system can recover what was intended, what was confirmed, what ran, what changed, and what evidence supports the result.

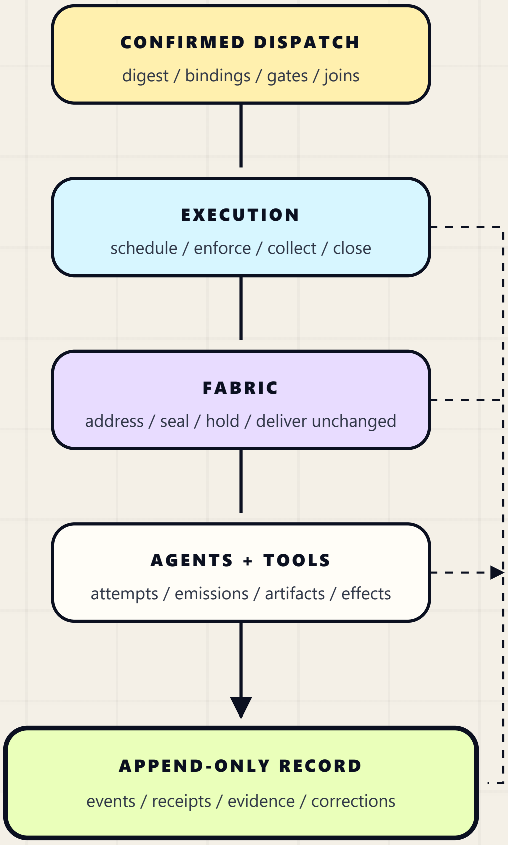

A trustworthy run preserves:

- **Stable identity.** Names and locations may change without breaking lineage.
- **Comparable shape.** Runs of one kind retain a basis for comparison.
- **Shape fidelity.** The confirmed topology is the topology executed.
- **Independent judgment.** Sealed work stays hidden until its release condition.
- **Unchanged delivery.** The recipient reads what the sender authored.
- **Recoverable acts.** Even empty-payload actions remain part of history.
- **Typed relations.** Every connection carries an explicit meaning.

Single-writer rule. Every service may emit facts; only Record writes the durable account. Corrections append instead of rewriting history.

Open boundary. The hypothesis records authority claims, but the source still leaves the authority provider and its verifiable proof contract unresolved.

> The test is simple: can a local effect explain the objective it served - and can the objective show how it became real?